

Cake classification

Machine Learning

Claudio Cusano

A.A. 2022/2023

Image recognition is a very important application of machine learning. There are two main approaches:

- with hand-crafted image features processed by a classification model;
- with deep-learning, with a neural network that directly processes the image pixels. Convolutional neural networks (CNN) are the most suitable architecture for this application.

1 Lab activity

In this exercise we will build classifiers for the classification of images of cakes. Before starting the exercise be sure to read the “Preliminaries” section here below, download the data set and review the relevant topics.


1.0 Preliminaries

The data set contains 120 images for each of 15 kinds of cake. For each class 100 images are in the training set and 20 form the test set. All the images have been resized to 224×224 pixels.

For feature extraction we will use low-level features. Some of them have been implemented in the `image_features.py` file. We will also use neural features, computed py using the ‘PVMLNet’ CNN.

1.0.1 PVMLNet

This CNN has been designed as a slight simplification over the AlexNet architecture. It takes as input color images of 224×224 pixels. A sequence of eight “same size” convolutions followed by three fully-connected layers (also implemented as convolutions) that compute the class scores as follows:



All layers are followed by ReLU activations. Some convolutions (S) are strided. A final softmax computes the class probabilities.

The network has been trained on the ILSVRC-12 subset of ImageNet. It achieves a top-1 accuracy of 61.0% and a top-5 accuracy of 83.3%.

1.1 Low-level features

Write a script that computes one of the low-level feature vector implemented in the `image_features.py` file. Train a classifier and evaluate the test accuracy.

1.2 Neural features

Use the pretrained PVMLNet to extract as features the activations of the last hidden layer. Train a perceptron without hidden layers and evaluate the test accuracy.

1.3 Transfer learning

Build a new network by replacing the last layer of PVMLNet with the weights of the trained perceptron.

2 Assignment

As homework, review and refine the scripts programmed in the lab activity. In addition, perform the following exercises.

2.1 Combining features

Try different combinations of low-level features (concatenate two or more feature vectors with `np.concatenate`).

2.2 Analysis

Identify the pairs of classes that are more likely to be confused with neural features. Also identify the test images that are misclassified even if the classifier predicted their label with high confidence.

2.3 Neural features

Try to use neural features computed by different hidden layers. When the activations are spatially distributed, you may reduce them to a single feature vector by averaging over the spatial dimensions.

2.4 Fine tuning (Optional)

Try to fine-tune the CNN after transfer learning (i.e. continue training with as very small learning rate like 10^{-5}).

2.5 Report

Prepare a report of one or two pages with the answers to document all the experiments and their results. The report must be in the PDF format. Include your name in the report and conclude the document with the following statement: “I affirm that this report is the result of my own work and that I did not share any part of it with anyone else except the teacher.”

Make a ZIP archive with the report and the python scripts (**do not include data, images, features etc.**) and upload it on the course website.